

Traffic Scheduling

Jeroen van Riel

November 2023

comments are in blue italics

1 Preliminaries

This section discusses the task of assigning arriving (platoons of) vehicles to time slots on a single intersection, which already gives rise to an interesting combinatorial optimization problem.

1.1 Single Machine Scheduling

Suppose we have n jobs that need processing on a single machine. The time required for each job $j \in \{1, \dots, n\}$ is called the *processing time* p_j . Once started, jobs may not be *preempted*. In order to obtain a valid *schedule*, we need to determine the start time y_j for each job j , such that at most one job is processing on the machine at all times. Let $C_j = y_j + p_j$ denote the *completion time* of job j . Our objective is to minimize the *total completion time* $C_1 + \dots + C_n$. By means of an *interchange argument*, it can be easily shown that an optimal solution is given by sorting the jobs according to increasing processing times, which is known as the Shortest Processing Time first (SPT) rule.

Now suppose that job j becomes available at its *release date* r_j , then a valid schedule requires $y_j \geq r_j$. It is not difficult to see that distinct release dates may require us to introduce *idle time* to obtain an optimal schedule. For example, consider $n = 2$ jobs with processing times $p_1 = 2, p_2 = 1$ and release dates $r_1 = 0, r_2 = \epsilon > 0$. Processing job 1 before job 2 results in a schedule with total completion time $\sum C_j = 2 + 3 = 5$. However, scheduling job 2 first requires us to introduce ϵ idle time, but has a better total completion time of $\sum C_j = (\epsilon + 1) + (\epsilon + 1 + 2) = 4 + 2\epsilon$. Schedules without idle time are called *non-delay*.

Next, we consider *precedence constraints* between jobs. Suppose that job j needs to be processed before job l , denoted as $j \rightarrow l$, then we simply require that $y_l \geq C_j$ in any feasible schedule. In particular, we may consider *chains* of precedence constraints. Let J_1, \dots, J_k be a partition of jobs into k non-empty families. For each family, we require that their jobs $J_l = \{j_1, \dots, j_{n_l}\}$ are

processed in the order $j_1 \rightarrow j_2 \rightarrow \dots \rightarrow j_{n_i}$, without loss of generality. Note that the order between jobs from different families is unspecified, which means that chains may be *merged* arbitrarily.

When job j is directly followed by job l , we might want to introduce *sequence-dependent setup time* s_{jl} , by requiring that $y_l \geq C_j + s_{jl}$. For our purposes, we will only consider setup times depending on the family to which jobs belong. Formally, for any pair of jobs $j_1 \in J_{l_1}$ and $j_2 \in J_{l_2}$ belonging to distinct families, we require that either

$$y_{j_2} \geq C_{j_1} + s_{l_1, l_2},$$

is satisfied or

$$y_{j_1} \geq C_{j_2} + s_{l_2, l_1}.$$

1.2 Single Intersection Scheduling

We are interested in the task of assigning arriving vehicles to a time slot on a single intersection. Using the modelling ingredients introduced above, we can now pose this task as a scheduling problem. Let vehicles be represented by jobs. We will refer to y_j as the *crossing time*, since it represents the time at which the vehicle starts crossing the intersection, which is modelled by the single machine. The time it takes before the next vehicle can enter the intersection is modelled by the processing time p_j . The earliest possible crossing time of a vehicle is modelled by its release date r_j . All vehicles that arrive to the intersection from the same lane belong to the same family. We assume that vehicles are driving on single-lane roads, which means that we do not allow *overtaking*. We model this by introducing chain precedence constraints based on the order of arrival (release dates). To guarantee safety, we require a setup time s_{l_1, l_2} between vehicles coming from different lanes. As the optimization objective, we consider the total completion time, since this is equivalent to minimizing the total delay experienced by all vehicles.

The problem of minimizing the total completion time with release dates (written as $1|r_j|\sum C_j$ in the three-field notation [1]) has been long known to be NP-hard [2], so there is no hope in finding a polynomial algorithm for our more general problem unless $\mathcal{P} = \mathcal{NP}$. Therefore, we must resort to exhaustive branch-and-bound methods or heuristics. However, note that this only holds for the problem in its full generality. It might, for example, be worthwhile to consider the complexity when one or more parameters are fixed. Specifically, it would be interesting to verify the complexity in case of a fixed number (e.g., two) of families.

Assumption 1.1. *The processing times $p_j = p$ and setup times $s_{l_1, l_2}s$ are all identical, for $j \in \{1, \dots, n\}$ and any pair of lanes l_1 and l_2 .*

For the moment, this leaves us with an accurate enough model for vehicles that cross the intersection without turning. For each job family $J = \{1, \dots, n\}$,




Figure 1: Illustration of the two possible sequences in Example 1.2.

observe that we may assume that $r_i = \max\{r_i, r_{i-1} + p\}$, without loss of generality, because of the chain precedence constraints. Before discussing solution methods, let us first consider some simple examples.

Example 1.1. Consider the situation with $J_1 = \{1\}, J_2 = \{2\}$. When both vehicles have the same release dates $r_1 = r_2 = r$, the order in which they cross the intersection does not influence the optimal total completion time of $\sum C_j = p + (p + S + p) = 3p + S$. Now, assume that vehicle 1 has a earlier release date, then it is easily seen that an optimal schedule requires vehicle 1 to go first.

Example 1.2. Consider the situation with $J_1 = \{1\}, J_2 = \{2, 3\}$. We are interested in how the release dates influence the order of the jobs in an optimal schedule. We assume that $r_3 = r_2 + p$. Suppose $r_1 = r_2$, then we see that 2, 3, 1 is optimal, which resembles some sort of “longest chain first” rule ([relate this to Algorithm 3.1.4 of Pinedo](#)). Now assume that $r_2 < r_1 + p + s$, otherwise 1, 2, 3 is simply optimal, because there is no *conflict* between the two lanes ([we will make this more precise later](#)). Furthermore, set $r_1 = 0$, without loss of generality. We compare the sequence 1, 2, 3 with 2, 3, 1. The first has optimal $\sum C_j = p + (p + s + p) + (p + s + p + p) = 6p + 2s$, while the second sequence has optimal $\sum C_j = (r_2 + p) + (r_2 + p + p) + (r_2 + p + p + s + p) = 3r_2 + 6p + s$. Therefore, we conclude that the second sequence is optimal if and only if

$$r_2 \leq s/3,$$

which roughly means that the “longest chain first” rule becomes optimal whenever the release dates are “close enough”.

Definition 1.1. A sequence of consecutive jobs $n, n+1, \dots, n+m \in F_l$ from one lane is called a *platoon* if and only if $r_{i+1} = r_i + p$ for $n \leq i < n+m$.

Example 1.3. Suppose we have exactly two platoons of vehicles $J_A = \{1, \dots, n_A\}$, $J_B = \{n_A+1, \dots, n_A+n_B\}$. Assuming that $n_A < n_B$ and $r_1 \leq r_{n_A+1}$, consider the ways the two platoons can merge by splitting A. Let k denote the number of vehicles of platoon A that go after platoon B and let $\sum C_j(k)$ denote the resulting total completion time. We have

$$\sum C_j(0) = \sum_{i=1}^{n_A+n_B} ip + n_B s,$$




Figure 2: Illustration of splitting platoon A in Example 1.3.

when A goes first and

$$\sum C_j(n_A + n_B) = \sum_{i=1}^{n_A+n_B} ip + n_A s,$$

when B goes first. For all the other cases, we have

$$\sum C_j(k) = \sum_{i=1}^{n_A+n_B} ip + n_B s + 2ks.$$

From this analysis, we observe that it is not beneficial to split platoon A under the current assumptions. Moreover, we see that platoon B should go first, whenever

$$r_{n_A+1} \leq (n_B - n_A)s / (n_A + n_B).$$

Definition 1.2. *As the examples above show, the computation of the total completion time always includes a common term*

$$\sum C_j = \sum_{i=1}^n ip + \dots,$$

which happens because of identical processing times (Assumption 1.1). Therefore, we will from now on discard this term while comparing different schedules, by defining the total setup delay

$$D = \sum C_j - \sum_{i=1}^n ip.$$

1.3 Single Platoon Positioning

Focus on the problem with two lanes F_A and F_B where we assume that there is only a single platoon arriving on lane B, i.e., the *single platoon positioning*

problem. In this case, the optimal schedule can be easily computed by considering all positions. However, we use this simple case to obtain some insight into when vehicles that are close act as if they are platoons.

Theorem 1.1. *Under Assumption 1.1, every vehicle scheduling problem must have an optimal schedule y in which every platoon $p = (n, n+1, \dots, n+m)$ satisfies $y_{j+1} = y_j + p$ for all $n \leq j < n+m$. So this is the $x_i = 0$ case.*

Proof. Assume there exists an optimal schedule y in which some platoon $P = (n, n+1, \dots, n+m)$ is *split* in two parts, i.e., there exists some $j > 0$ such that $y_{n+j} > y_{n+j-1} + p$. Whenever the platoon is split in more parts, the following argument can be applied exhaustively.

instead of considering parts P_1 and P_2 , just show how a single vehicle can be scheduled earlier and argue that this argument may be applied exhaustively

Let P_1 and P_2 denote the separate parts of the original platoon. Let B denote the set of vehicles that have $C_{n+j-1} + s \leq y_l \leq y_{n+j} - s$, i.e., those that are scheduled between the two parts of P . Similarly, let A be the set of vehicles that have $y_l \geq C_{n+m} + s$, i.e., those that are scheduled after P_2 .

Let x be the time between the end of P_1 and the start of the first job in B . The time between the last job of B and the first job of P_2 must be s , otherwise the total completion time can be simply decreased by shifting P_2 earlier in time. Let y be the time between the end of P_2 and the start of the first job in A .

Now put P_2 right after the end of P_1 , such that $y_{n+j} = y_{n+j-1} + p$ and consider the resulting change in total completion time.

explain our method of idle time times number of jobs method □

Assume that the first vehicle on the intersection arrives from lane A . We provide a criterium for vehicles from A to be candidates for going before the B -platoon, with length (number of vehicles) n_B .

Proposition 1.1. *Assume there is a vehicle $0 \in F_A$ starting at $y_0 = 0$. Let x_i denote the finish-start delay between vehicle $i-1$ and i when scheduled as early as possible, so let*

$$x_i = r_i - (r_{i-1} + p).$$

If there exists a smallest integer i' such that $x_{i'} \geq 2s + n_B p$, then any vehicle $i \geq i'$ needs to go after B in any optimal schedule.

illustrate the two job case with the “decision” diagram

1.4 Branch-and-Bound

Formulate the problem as a MIP and solve using branch-and-bound. Argue that imposing maximum vehicle delay allows us to limit the set of disjunctions that we need to consider.

lazy row generation
define conflicts

1.5 Insertion Heuristic

We propose a heuristic that constructs a schedule for two lanes by considering the vehicles in order of arrival and iteratively inserting each next vehicle in the current partial schedule.

As in the single platoon scheduling problem, let F_A and F_B denote two lanes, but now we allow lane B to have arbitrary arriving vehicles. For the first platoon of B , find the single platoon position among vehicles of A . Then for every next platoon of B , solve the single platoon positioning problem, while treating the position of the previously scheduled platoons fixed.

The following example shows that the resulting schedule may not be optimal, because the position of earlier scheduled platoons restricts the possible positions of later platoons.

Example 1.4. *come up with example*

2 Job Shop Scheduling

This section discusses the classical *job shop* scheduling environment. Assume there are m machines and n jobs. A job j consists of exactly m operations, one for each of the machines. We let (i, j) denote the operation of job j that needs to be processed on machine i . The time required for processing operation (i, j) is denoted by p_{ij} . The operations of a job need to be executed in a fixed given order, which may be different among jobs, and an operation may only start once its predecessor has completed processing. Each machine can process at most one operation at the same time and, once started, operations cannot be preempted.

Let the set of all operations be denoted by N . Furthermore, we encode the job routes by defining the set \mathcal{C} of precedence constraints $(i, j) \rightarrow (k, j)$. Let y_{ij} denote the start of operation (i, j) . The completion time of job j is defined as $C_j := y_{lj} + p_{lj}$, where l is the last machine on which j must be processed. A valid schedule is given by setting values for y_{ij} such that the above requirements are met. There are various measures for how *good* a given schedule is. For the purpose of this example, let us consider the well-known makespan objective $C_{\max} := \max_j C_j$, which is often related to efficient use of the available machines. Minimizing the makespan can now be formulated as a Mixed-Integer Linear Program (MILP) as follows:

$$\begin{aligned}
& \text{minimize } C_{\max} \\
& y_{ij} + p_{ij} \leq y_{kj} && \text{for all } (i, j) \rightarrow (k, j) \in \mathcal{C} \\
& y_{il} + p_{il} \leq y_{ij} \text{ or } y_{ij} + p_{ij} \leq y_{il} && \text{for all } (i, l) \text{ and } (i, j), i = 1, \dots, m \\
& y_{ij} + p_{ij} \leq C_{\max} && \text{for all } (i, j) \in N \\
& y_{ij} \geq 0 && \text{for all } (i, j) \in N
\end{aligned}$$

The first set of constraints enforce the order of operations belonging to the same job. The second set of constraints are called *disjunctive*, because they model that we need to choose between jobs j and l to be scheduled first on machine i . The third set of constraints are used to define the makespan and the last line enforces non-negative start times.

A common way of representing job shop problems is through the corresponding *disjunctive graph* ([cite! is there a more accessible reference than this original french paper?](#)), see Figure 3 for an example. Let $(N, \mathcal{C}, \mathcal{D})$ be a directed graph with nodes corresponding to all the operations and two types of arcs. The *conjunctive* arcs \mathcal{C} encode the given machine order for each job. If $(i, j) \rightarrow (k, j) \in \mathcal{C}$ then job j needs to be processed on machine i before it may be processed on machine k . The *disjunctive* arcs \mathcal{D} are used to encode the decisions regarding the ordering of jobs on a particular machine. For every pair of different jobs j and l that need to be processed on the same machine i , there is a pair of arcs in opposite directions between (i, j) and (i, l) . Therefore, each machine corresponds to a clique of double arcs between all the operations that need to be performed on that machine.




Figure 3: Disjunctive graph corresponding to a job shop instance with three jobs and three machines. Operations that need to be scheduled on the same machine have been given the same color, emphasizing the cliques formed by their disjunctive arcs (drawn as dashed lines).

For each machine k , we need to order the operations that need processing on k . This corresponds to choosing exactly one arc from every pair of opposite disjunctive arcs. Let $\mathcal{D}' \subset \mathcal{D}$ be such a selection of disjunctive arcs and let $G(\mathcal{D}')$ denote the resulting graph, to which we will refer as a *complete* disjunctive graph. It is not hard to see that any feasible schedule corresponds to an acyclic complete disjunctive graph.

2.1 Schedule generation

There are some classes of schedules that are relevant for job shop problems. The following definitions are taken from the textbook by Pinedo [3].

Definition 2.1 (Non-Delay Schedule). *A feasible schedule is called non-delay if no machine is kept idle while an operation is waiting for processing.*

Definition 2.2 (Active Schedule). *A feasible non-preemptive schedule is called active if it is not possible to construct another schedule, through changes in the order of processing on the machines, with at least one operation finishing earlier and no operation finishing later.*

Definition 2.3 (Semi-Active Schedule). *A feasible non-preemptive schedule is called semi-active if no operation can be completed earlier without changing the order of processing on any one of the machines.*

It is clear that active schedules are necessarily semi-active, but the converse does not always hold. Furthermore, when preemption is not allowed, non-delay




Figure 4: Optimal schedule for Example 2.1. The dotted lines indicate unit time steps. Machine 2 is kept idle, while operation (2,2) could have already been scheduled at time 1.

schedules are necessarily active. For some scheduling problems, it can be shown that an optimal schedule exists that belongs to such classes.

Example 2.1. Consider two jobs on three machines with processing times $p_{\cdot 1} = (2 \ 2 \ 4)^T$ and $p_{\cdot 2} = (1 \ 2 \ 1)^T$. It is easily seen that the makespan is determined by $C_{\max} = p_{11} + p_{21} + p_{31} = 8$ and an optimal schedule is given in Figure 2.1. This schedule is not non-delay, because machine 2 is kept idle from time 2 to 3, but operation (2,2) could be scheduled here.

For job shop problems, the optimal schedule is not necessarily non-delay, as the example illustrates for the makespan objective. However, it can be shown that $Jm||\gamma$ has an active optimal schedule, when γ is a regular objective, which means that it is a non-decreasing function of the completion times C_1, C_2, \dots, C_n . The makespan and total completion time are examples of regular objectives. Because of this result, branch & bound methods have been proposed based on enumeration of all active schedules.

also active optimal schedule for $Jm|r_j|\gamma$?

Proposition 2.1. *Every complete disjunctive graph corresponds to a unique semi-active schedule.*

Proof. Given some semi-active schedule, it is obvious that the corresponding disjunctive graph can be constructed by simply encoding the ordering of the operations in the schedule by the disjunctive arcs.

We now argue that every complete disjunctive graph cannot correspond to more than one semi-active schedule. Recursively compute a lower bound on the starting time $LB(O_{jk}) = \max_{n \in O_{jk}^-} LB(n) + p_{jk}$ where O_{jk}^- denotes the set of in-neighbors. In other words, $LB(O_{jk})$ is the longest path, when we associate weight p_{jk} with every incoming arc of node O_{jk} and add a dummy node s with arcs to all first operations. Every feasible schedule y must satisfy $y_{jk} \geq LB(O_{jk})$. Hence, if any $y_{jk} > LB(O_{jk})$, then it is not semi-active. \square

Remark 2.1. *After adding another dummy sink node t and adding arcs with weight zero from all the last operations to it, it is clear that $LB(t) = C_{\max}$.*




Figure 5: Disjunctive graph corresponding to Example 2.1 augmented with source, sink and weights. The longest path determining the makespan $C_{\max} = LB(t)$ is highlighted.

show how to adapt for total completion time (see Pinedo 7.3)

Include here our solution of Exercise 7.11 from Pinedo, i.e., a branching scheme for generating all schedules based on inserting disjunctive arcs. From the resulting disjunctive graph, all semi-active schedules follow by applying the machine makespan rule, introduced in the next section.

2.2 Priority dispatching

A *priority dispatching rule* selects a job from a list of currently available (unscheduled) jobs to be scheduled next. This produces a sequence χ of jobs that are placed in a *partial schedule*. Manually crafted priority dispatching rules are mostly based on elementary features of the current partial schedule. Well-known examples include the job with the most/least work remaining (MWR/LWR) or the job with the shortest/longest processing time (SPT/LPT).

Determining the dispatching sequence alone is not enough to obtain a schedule y . The exact starting times are determined by a *placement rule*. A sensible placement rule would be to put the job in the earliest gap of idle time where this job fits, i.e., where it can be placed without violating any constraints. We call this the *earliest gap* rule. Alternatively, consider the straightforward placement rule that chooses the smallest starting time y_{ij} that satisfies $y_{ij} \geq y_{il} + p_{il}$ for all jobs l that have already been placed on machine i in the current partial schedule and $y_{ij} \geq y_{kj} + p_{kj}$, when (k, j) is the preceding operation. We call this the *last position* rule, because the job is placed as the last job on the machine.

We are now interested in whether scheduling with dispatching rules can generate all necessary schedules. More precisely, we wonder whether all active schedules can be generated from dispatching.

Proposition 2.2. *For every feasible semi-active schedule, there exists a sequence χ that generates it using the last position rule.*

Proof. Consider an arbitrary semi-active schedule y and let G be the corresponding complete disjunctive graph, which is unique by Proposition 2.1 and has $y_{jk} = LB(O_{jk})$ for each node $O_{jk} \in N$. Because G must be acyclic, there exists some topological ordering χ of the nodes. Therefore, dispatching in order of χ with the machine makespan rule will produce exactly the schedule with $y_{jk} = LB(O_{jk})$. \square

3 Job Shop with Reinforcement Learning

3.1 Zhang et al.

The authors of [4] use reinforcement learning to solve the job shop problem with makespan objective $Jm||C_{\max}$. Specifically, the method learns dispatching rules, so actions correspond to selecting which (unfinished) job to schedule next. States are represented by a disjunctive graph corresponding to the partial schedule. The policy is parameterized by a graph neural network that reads the current disjunctive graph. They use a dense reward based on the current lower bound of the makespan.

The state transition example shown in Figure 2 of [4] suggests yet another placement rule, because operation O_{32} is placed before O_{22} , but it does not fit the gap. Consequently, operation O_{22} has to be delayed by three time steps.

Formalize the placement rule that Zhang et al. use and argue that this does not guarantee that all active schedules can be generated in this way.

4 Traffic Scheduling in Networks

4.1 MIP Formulation

We now turn to vehicles traveling through a network of intersections. The network may be thought of as a weighted directed graph $G = (V, E)$, with nodes and arcs representing intersections and roads, respectively. **For now, we assume that the graph is acyclic.** Let $d(x, y)$ be defined as the *distance* between nodes x and y . We assume there are no nodes of degree two, since their two incident arcs (x, y) and (y, z) could be merged into one arc (x, z) with $d(x, z) = d(x, y) + d(y, z)$, without loss of expressiveness. Furthermore, we assume the graph is connected. Each node of degree one is called an *external node* and models the location where vehicles enter (*entrypoint*) or exit (*exitpoint*) the network. A node of degree at least three is called an *internal node* and models an intersection.

Each vehicle j enters the network at some external node s and follows a predetermined sequence of arcs $R = ((s, i_1), (i_1, i_2), \dots, (i_{n-1}, i_n), (i_n, d))$ towards an external node d where it leaves the network. Vehicles are not able to overtake each other when traveling on the same arc. We assume that arcs provide infinite *buffers* for vehicles, meaning that there is no limit on the number of vehicles that are traveling on the same arc at the same time. However, we impose a minimum time required to travel along an arc (x, y) . By assuming identical maximum speeds among vehicles, $d(x, y)$ can be directly interpreted as this minimum travel time.

Let $y(i, j)$ denote the time vehicle j enters intersection i . Crossing an intersection takes p time per vehicle and at most one vehicle can cross an intersection at the same time. When two consecutive vehicles crossing an intersection originate from the same arc, they may pass immediately after each other. However, when a vehicle j_1 that wants to cross comes from a different arc than the vehicle j_0 that last crossed the intersection, we require that there is at least a *switch-over time* S between the moment j_0 leaves and the moment j_1 enters the intersection.

Assuming that arrival times and routes of all vehicles are fixed and given, our task is to determine when individual vehicles should cross intersections by setting values for $y(i, j)$. This problem is similar to job shop scheduling, with intersections and vehicles now taking the roles of machines and jobs, and can also be formulated as a MIP, as we will show below.

The main difference with job shop scheduling is related to the ordering of vehicles. In the job shop model, every order $o \in \sigma_n$ of jobs on a machine i was allowed. By assuming that vehicles cannot overtake each other, we are limiting the valid orderings. Given two routes R_j and R_l , we define a common path $p = (i_1, \dots, i_L)$ as a substring of both routes. We refer to the first node i_1 of a common path as a *merge point*. The set of all common paths is denoted by P_{jl} .

Consider a common path p whose merge point i_1 is an external node. In this case, i_1 must be the entrypoint of both vehicles. That means that the order on p is determined by the order of arrival. Let r_j denote the arrival time (called

the *release date* in scheduling) of vehicle j . If $r_j < r_l$, then we require that j goes first, formally stated as

$$y_{ij} + p_{ij} \leq y_{il} \text{ for all } i \in p.$$

Now suppose that j and l have a mergepoint $i_1 \in p$ that is not an external node, then these vehicles approach i_1 from different arcs, which means that we have what we will call a *conflict* between j and l , because the scheduler must choose which vehicle crosses the intersection first. Furthermore, the switch-over time should be respected at i_0 . Requiring that vehicle j must cross first can be formally stated as

$$y_{i_1j} + p_{i_1j} + S \leq y_{i_1l} \text{ and } y_{ij} + p_{ij} \leq y_{il} \text{ for all } i_1 \neq i \in p.$$

In order to model the above decision making with a MIP, we introduce binary decision variables $s_{i,j,l}$ to encode the relative order of j and l at some common node $i \in p \in P_{jl}$. A value of zero corresponds to j crossing first and a value of one indicates that l crosses first. In the case of i_1 being an external node, we treat $s_{i,j,l}$ as a fixed parameter of the MIP. For each common path $p \in P_{jl}$, we must have

$$s_{i_1,j,l} = s_{i_2,j,l} = \dots = s_{i_L,j,l}.$$

Let N denote the set of operations like in the job shop example, but now pairs correspond to vehicles and the intersections they encounter along their route. Also similarly, let A encode the route constraints of each vehicle. We use p_1 to denote the merge point of path p . Let P_{jl}^e denote all the common paths that have an external merge point and let P_{jl} denote the other common paths. By implicitly letting the indices j and l run over all the *ordered pairs* of vehicles, we write

$$\text{maximize } P(y) \tag{1a}$$

$$y_{ij} + p_{ij} + t_{ik} \leq y_{kj}, \quad \text{for all } (i, j) \rightarrow (k, j) \in A, \tag{1b}$$

$$y_{ij} + p_{ij} \leq y_{il}, \text{ for all } i \in p, \quad \text{for all } p \in P_{jl}^e, s_{p_1jl} = 1, \tag{1c}$$

$$y_{p_1j} + p_{p_1j} + S \leq y_{p_1l}, \quad \text{for all } p \in P_{jl}, s_{p_1jl} = 1, \tag{1d}$$

$$y_{ij} + p_{ij} \leq y_{il}, \text{ for all } i \neq p_1, \quad \text{for all } p \in P_{jl}, s_{p_1jl} = 1, \tag{1e}$$

$$y_{ij} \geq 0, \quad \text{for all } (i, j) \in N, \tag{1f}$$

where $P(y)$ denotes some unspecified performance metric, leaving for now the question of what it means for a schedule to be *good*. In order to solve this program using existing solvers, we use the well-known *big-M* method to encode which constraints from sets (1c), (1d) and (1e) are active for specific values of the binary variables s .

4.2 Structure in Network Schedules

Before we start experimenting with the above problem (see if the MIP method scales and later trying to learn policies with (MA)RL), we first should try to

find some common structural patterns in solutions or even some simple general scheduling rules.

- Study simple situations with the total completion time objective ($\sum C_j$), in which some kind of LPT rule holds. (*explain why LPT-like schedules arise*)
- Study tandem networks where the $\sum C_j$ objective causes platoon splitting due to some kind of “propagation of delay costs”. (*insert the example that I have*)
- Can we study this “delay cost propagation” using some kind of dependency graph, e.g., by recording that delaying job A would require us to also delay job B, which would also require to delay C, etc.?

4.3 Platoon Splitting

In the single intersection case, platoon splitting is never necessary. So in this context, platoon splitting can only become necessary from using a different performance metric, e.g., one that takes into account fairness. Once we start looking at the network-level interactions, however, we can show that platoon splitting is sometimes necessary to obtain an optimal schedule. The main idea is that delaying a job j might require to also delay further downstream jobs. Therefore, it might be cheaper overall to not delay j , but interrupt some already running job l instead.

References

- [1] R. L. Graham, E. L. Lawler, J. K. Lenstra, and A. H. G. R. Kan, “Optimization and Approximation in Deterministic Sequencing and Scheduling: A Survey,” in *Annals of Discrete Mathematics* (P. L. Hammer, E. L. Johnson, and B. H. Korte, eds.), vol. 5 of *Discrete Optimization II*, pp. 287–326, Elsevier, Jan. 1979.
- [2] J. K. Lenstra, A. H. G. Rinnooy Kan, and P. Brucker, “Complexity of Machine Scheduling Problems,” in *Annals of Discrete Mathematics* (P. L. Hammer, E. L. Johnson, B. H. Korte, and G. L. Nemhauser, eds.), vol. 1 of *Studies in Integer Programming*, pp. 343–362, Elsevier, Jan. 1977.
- [3] M. L. Pinedo, *Scheduling: Theory, Algorithms, and Systems*. Cham: Springer International Publishing, 2016.
- [4] C. Zhang, W. Song, Z. Cao, J. Zhang, P. S. Tan, and C. Xu, “Learning to Dispatch for Job Shop Scheduling via Deep Reinforcement Learning,” Oct. 2020.