

Trajectory Optimization of Autonomous Vehicles in Networks of Intersections

Jeroen van Riel

March 2025

Contents

1	Trajectories in tandem of two intersections	1
1.1	Single vehicle	2
1.2	Multiple vehicles	4
1.2.1	Ad-hoc approach	5
1.2.2	Schedule time approach	6
2	Trajectories in networks	7
2.1	General decomposition	8
2.2	Decomposition for delay objective	8
2.3	Crossing time scheduling	9
3	Constructive heuristics	9
3.1	Exhaustive heuristic	10
3.2	Neural constructive heuristic	10
3.3	Network-agnostic architecture	11

1 Trajectories in tandem of two intersections

When considering multiple vehicle driving between intersection, we can no longer ignore the issue of road capacity, because the fact that only a limited number of vehicles can drive or wait at the same time on a lane between intersections may cause network-wide effects. The capacity of lanes between intersections is intimately related to the trajectories of vehicles, which we first want to understand better. We have been using an optimal control formulation




Figure 1: Tandem of two intersections v and w with lane of length $d(v, w)$. The grey rectangle represents some vehicle that just left intersection v . We will assume that a vehicle has maximum speed as long as it occupies an intersection, so it is only allowed to decelerate once it passes this position.

with the objective that keeps the vehicles as close as possible to the next intersection at all times (**MotionSynthesize**). This problem can be solved using direct transcription, which works well enough if we just want to simulate the behavior of the system. However, we believe that it is possible to explicitly formulate the optimal controller. We will explain how to compute trajectories corresponding to those obtained by direct transcription, but without using time discretization.

Before we turn to the general case of networks of intersection, we will first investigate the trajectories of vehicles in a tandem of two intersections as depicted in Figure 1. Let v denote the left intersection and w the right intersection and assume that vehicles drive from left to right. Furthermore, we will call the road segment strictly between both intersection areas the *lane*. To facilitate the following discussion, we will use p to denote the position of a vehicle to distinguish it from the state $x = (p, v)$, which also includes the velocity, and we fix position $p = 0$ at the stop-line of intersection w . Let the length and width of a vehicle i be denoted by L_i and W_i , respectively. We measure the position of a vehicle at the front bumper. We will make the following assumption, that allow us to easily derive explicit expressions of these trajectories.

Assumption 1.1. *All vehicles have the same length $L_i = L$ and width $W_i = W$. Lanes are exactly W units wide and are axis-aligned, such that intersection are squares.*

Assumption 1.2. *Vehicles must drive at full speed when entering an intersection and keep driving at full speed as long as they occupy an intersection.*

Now assume that some vehicle is scheduled to exit v at time t_0 and to enter w at some time t_f . Let $p_0 = -d(v, w) + W$ denote the position of the vehicle when it starts to exit v . Let $y(t)$ denote the trajectory of the vehicle that drives in front of the current one, assuming it exists. In order to keep the vehicle as close to w as possible at every time, while respecting the double integrator vehicle dynamics $\dot{p} = v$, $\ddot{p} = u$, we can generate a trajectory by solving the optimal control problem

$$\begin{aligned} \max_u \quad & \int_{t_0}^{t_f} p(t) dt \\ \text{s.t.} \quad & 0 \leq v(t) \leq v_{\max}, \\ & -a_{\max} \leq u(t) \leq a_{\max}, \\ & y(t) \leq p(t), \\ & p(t_0) = p_0, \quad p(t_f) = 0, \\ & v(t_0) = v_{\max}, \quad v(t_f) = v_{\max}. \end{aligned} \tag{1}$$

This problem can be solved by using a direct transcription method. After observing some example solutions, we believe that the optimal control should switch between no acceleration and either full acceleration or full deceleration, i.e., we have a control function $u(t) := \ddot{x}(t)$ satisfies $u(t) \in \{-a_{\max}, 0, a_{\max}\}$ and some sequence of alternating deceleration and acceleration periods, represented by some sequence of disjoint intervals

$$(D_0, A_1, D_1, \dots, A_{n-1}, D_{n-1}, A_n),$$

so that the optimal controller is given by

$$u(t) = \begin{cases} -a_{\max} & \text{if } t \in D_k \text{ for some } k, \\ a_{\max} & \text{if } t \in A_k \text{ for some } k, \\ 0 & \text{otherwise.} \end{cases}$$

1.1 Single vehicle

We can proof the above hypothesis and provide explicit expressions for D_k and A_k of the optimal trajectory whenever the safe following constraint $y(t) \leq p(t)$ is never active and

can essentially be ignored, either because there is no vehicle ahead or it is sufficiently far away. The proof uses a variant of the Pontryagin Maximum Principle, specifically tailored to problems with state constraints [1].




Figure 2: Tandem of intersections with distances used in the determination of the lane capacity and the waiting positions.

1.2 Multiple vehicles

By observing optimal trajectories generated using direct transcription, we observe that there is a lot of structure in how the trajectory of a vehicle influences the trajectories of later vehicles. In this section, we provide a direct way to explicitly calculate trajectories, without proving that these trajectories are always feasible and optimal. Instead, we verify correctness of our method numerically by comparing the output to direct transcription for various system parameters and crossing times.

We start by investigating the limit on the number of vehicles that can occupy the lane at the same time. Imagine that vehicles enter the lane until it is full and then only after all vehicles have come to a full stop, they start to leave the lane by crossing w . Without any additional constraints, it follows from the vehicle dynamics that the time it takes to fully accelerate from rest to maximum velocity is given by $d_f = v_{\max}/a_{\max}$ and with corresponding trajectory

$$p^*(t) = a_{\max}t^2/2 \quad \text{for } 0 \leq t \leq d_f.$$

Because we assume that maximum deceleration is also a_{\max} , it also takes d_f time to fully decelerate from maximum velocity to rest. Suppose that we want to design the tandem network such that at least $c(v, w)$ vehicles can enter and decelerate to some waiting position, from which it is also possible to accelerate again to full speed before crossing w . Vehicles are required to drive at full speed as long as they occupy any intersection. Therefore, a vehicle crossing v can only start decelerating after $p(t) \geq W + L$, so the earliest position where a vehicle can come to a stop is $p = W + L + p^*(d_f)$. Because vehicles need to gain maximum speed before reaching w , the position closest to w where a vehicle can wait is $-p^*(d_f)$. Hence, in order to accomodate for $c(v, w)$ waiting vehicles, the length of the lane must satisfy

$$d(v, w) \geq W + L + 2p^*(d_f) + (c(v, w) - 1)L,$$

as illustrated in Figure 2. Conversely, given the lane length $d(v, w)$, the corresponding lane capacity is given by

$$c(v, w) = \text{floor} \left(\frac{d(v, w) - W - 2p^*(d_f)}{L} \right),$$

where $\text{floor}(x)$ denotes the largest integer smaller than or equal to x .

It turns out that the fixed locations where vehicles wait in the above scenario are helpful in describing the optimal trajectories, even when vehicles never fully stop. We will denote these fixed *waiting positions* as

$$p_k = -p^*(d_f) - (c(v, w) - k)L,$$

for $k = 1, \dots, c(v, w)$. Furthermore, let $p_d = p_1 - p^*(d_f)$ denote the position from which vehicles must decelerate in order to stop at the earliest waiting position p_1 , which is the farthest from the destination intersection. Now consider a vehicle that moves from p_k to the next waiting position p_{k+1} , so it moves exactly distance L . We consider such a start-stop movement, without considering any safe following constraints. By symmetry of the control constraints, the vehicle moves the same distance during both acceleration and deceleration. Furthermore, the vehicle needs to be at rest at the start and end of such trajectory. Hence, it is clear that it takes the same amount of time d_s to accelerate and decelerate. We assume that $d_s < d_f$, which ensures that v_{\max} is never reached during the start-stop movement, which is illustrated in Figure 3. In this case, it is clear that we must have $L = 2p^*(d_s)$, from which we derive that $d_s = \sqrt{L/a_{\max}}$.

1.2.1 Ad-hoc approach

We will now present a method to calculate the trajectory of a vehicle based on its crossing times at v and w and the trajectory of the vehicle ahead of it. We start from a sequence of deceleration and acceleration intervals that are possibly overlapping and then merge them in pairs from left to right until they become disjoint. Specifically, let

$$t_d := t_0 + \frac{d(v, w) - W - p_d}{v_{\max}}$$

be the start of the initial deceleration interval $D_0 := [t_d, t_d + d_f]$, which is exactly the moment the vehicle needs to start decelerating in order to stop at the first waiting position p_1 , so at time t_d the vehicle is at position p_d . Similarly, let $t_a = t_f - d_f$ be the start of the final acceleration interval $A_f := [t_a, t_a + d_f]$. Now for every $k = 1, \dots, c(v, w) - 1$, we consider a pair of start-stop intervals $S_k := (A_k, D_k) = ([t_k, t_k + d_s], [t_k + d_s, t_k + 2d_s])$ at some starting time t_k , moving the vehicle from p_k to p_{k+1} . We first show how to merge these intervals such that a sequence of disjoint intervals is obtained. Afterwards, we show how times t_k follow from the trajectory of the vehicle ahead.

First, we show how to merge D_0 and S_1 . When we have $t_d + d_f \leq t_1$, it is clear that both parts are already completely disjoint, so we do not have to do anything further. As t_d and t_1 get closer than d_f , we need to start merging D_0 and A_1 as we will show. In this case, we observe that D_0 gets shorter at the end by some ϵ and the acceleration part of S_1 gets shorter at the beginning, also by the same amount ϵ , because the velocities must match. More precisely, we construct two new consecutive intervals

$$\begin{aligned} D'_0 &= [t_1 + 2\epsilon - d_f, t_1 + \epsilon], \\ A'_1 &= [t_1 + \epsilon, t_1 + d_s], \\ D_1 &= [t_1 + d_s, t_1 + 2d_s]. \end{aligned}$$

We now determine ϵ as a function of $t_1 - t_d$. Because D'_0 and A'_1 are both ϵ shorter, the total distance that is traversed is now $2p^*(\epsilon)$ shorter. This means that the start of D'_0 should be $2p^*(\epsilon)/v_{\max}$ later than t_d , which gives equation

$$t_d + \frac{2p^*(\epsilon)}{v_{\max}} = t_1 + 2\epsilon - d_f,$$

which is quadratic in ϵ . Because we must have $\epsilon < d_f$, we need the smallest solution

$$\epsilon = d_f - \sqrt{d_f(t_1 - t_d)}.$$

When we keep increasing t_d while keeping t_1 fixed, we observe that eventually part A_1 will completely disappear and D_0 and D_1 will become a single interval, which is easily seen to happen when $\epsilon \geq d_s$. Equivalently, this happens when t_d is such that $t_d \geq t_1 - d_f + 2d_s - 2p^*(d_s)/v_{\max}$.




Figure 3: Shape of some start-stop trajectory of a single isolate vehicle moving forward from some current waiting position p_k to the next p_{k+1} .

We now show how two start-stop parts have to be merged if they overlap. Consider two start-stop parts A_k, D_k and A_{k+1}, D_{k+1} and suppose that $t_k + 2d_s < t_{k+1}$ such that both parts have to be merged. The parts D_k and A_{k+1} merge by removing ϵ on both sides, similarly as above. However, this causes the D_k, A_{k+1}, D_{k+1} parts to shift up. Therefore, A_k and D_k each need to be lengthened at the side where they meet by some δ to match this. Hence, it turns out we have to use the intervals

$$\begin{aligned} A'_k &= [t_k, t_k + d_s + \delta], \\ D'_k &= [t_k + d_s + \delta, t_{k+1} + \epsilon], \\ A'_{k+1} &= [t_{k+1} + \epsilon, t_{k+1} + d_s], \\ D_{k+1} &= [t_{k+1}, t_{k+1} + d_s]. \end{aligned}$$

We have that ϵ and δ need to satisfy

$$\begin{cases} 2\delta + 2\epsilon = t_{k+1} - t_k, \\ 2p^*(d_s + \delta) - 2p^*(d_s - \epsilon) = L. \end{cases}$$

Solving this system of equations yields

$$\begin{aligned} \delta &= \frac{L/a_{\max}}{t_{k+1} - t_k} - d_s + \frac{t_{k+1} - t_k}{4}, \\ \epsilon &= (t_{k+1} - t_k)/2 - \delta. \end{aligned}$$

Finally, we consider the merge with the final acceleration bang.

Note that the above types of merging are enough to process the whole sequence of bangs, because when the first merge of D_0 and A_1, D_1 results in a single D'_1 and there is a next A_2, D_2 , we are again in the first situation.

We now show how t_k follow from the trajectory of the preceding vehicle.

1.2.2 Schedule time approach




Figure 4: Sketch of vehicle trajectory in tandem with all the full start-stop parts unmerged.




Figure 5: Illustration of some grid-like network of intersections with vehicles drawn as grey rectangles. There are five vehicle routes: two from east to west and three from south to north. Turning at intersections is not allowed.

2 Trajectories in networks

We now extend the single intersection model to a network of intersections without turning routes, illustrated in Figure 5. We define a directed graph (V, E) with nodes V and arcs E , representing the possible paths that vehicles can follow. Nodes of in-degree at least two are called *intersections*. Nodes with only outgoing arcs are *entrypoints* and nodes with only incoming arcs are *exitpoints*. Let $d(v, w)$ denote the distance between nodes v and w . For each route index $r \in \mathcal{R}$, we let

$$\bar{V}_r = (v_r(0), v_r(1), \dots, v_r(m_r), v_r(m_{r+1}))$$

be the path that vehicles $i \in \mathcal{N}_r$ follow through the network. We require that the first node $v_r(0)$ is an entrypoint and that the last node $v_r(m_{r+1})$ is an exitpoint and we write

$$V_r = \bar{V}_r \setminus \{v_r(0), v_r(m_{r+1})\}$$

to denote the path restricted to intersections. We say that some $(v, w) \in E$ is on path V_r whenever v and w are two consecutive nodes on the path and we write E_r to denote the set of all these edges. We require that routes can only overlap at nodes by making the following assumption.

Assumption 2.1. *Every arc $(v, w) \in E$ is part of at most one route V_r .*

We start by considering networks in which all roads are axis-aligned such that intersections always involve perpendicular lanes and where routes are such that no turning is required. For each $v \in V_r$ define the conflict zone $\mathcal{E}_r(v) = (b_r(v), e_r(v))$ and consider the union

$$\mathcal{E}_r = \bigcup_{v \in V_r} \mathcal{E}_r(v)$$

corresponding to the positions of vehicles $i \in \mathcal{N}_r$ for which it occupies an intersection on its path V_r . By reading $\mathcal{E}_i \equiv \mathcal{E}_r$ for $r(i) = r$, the single intersection problem naturally extends to the network case. Like before, the resulting problem can be numerically solved by a direct transcription method.

2.1 General decomposition

The general two-stage decomposition for the single intersection extends rather naturally to the present model. Let for each pair (i, v) of some vehicle $i \in \mathcal{N}$ and an intersection $v \in V_{r(i)}$ along its route, let

$$\inf\{t : x_i(t) \in \mathcal{E}_r(v)\} \quad \text{and} \quad \sup\{t : x_i(t) \in \mathcal{E}_r(v)\}$$

be the crossing time and exit time, which we denote by $y(i, v)$ and $y(i, v) + \sigma(i, v)$, respectively. Instead of a single set of conflicts, we now define for each intersection $v \in V$ in the network the set of conflict pairs

$$\mathcal{D}^v = \{\{i, j\} \subset \mathcal{N} : r(i) \neq r(j), v \in V_{r(i)} \cap V_{r(j)}\}.$$

Now the two-stage approach is to solve

$$\begin{aligned} \min_{y, \sigma} \quad & \sum_{r \in \mathcal{R}} F(y_r, \sigma_r) \\ \text{s.t.} \quad & y(i, v) + \sigma(i, v) \leq y(j, v) \text{ or} \\ & y(j, v) + \sigma(j, v) \leq y(i, v), \quad \text{for all } \{i, j\} \in \mathcal{D}^v \text{ and } v \in V, \\ & (y_r, \sigma_r) \in \mathcal{S}_r, \quad \text{for all } r \in \mathcal{R}, \end{aligned}$$

where $F(y_r, \sigma_r)$ and \mathcal{S}_r are the value function and set of feasible parameters, respectively, of the parametric trajectory optimization problems

$$\begin{aligned} F(y_r, \sigma_r) = \min_{x_r} \quad & \sum_{i \in \mathcal{N}_r} J(x_i) \\ \text{s.t.} \quad & x_i(t) \in D_i(s_{i,0}), \quad \text{for } i \in \mathcal{N}_r, \\ & x_i(y(i, v)) = b_r(v), \quad \text{for } v \in V_r, i \in \mathcal{N}_r, \\ & x_i(y(i, v) + \sigma(i, v)) = e_r(v), \quad \text{for } v \in V_r, i \in \mathcal{N}_r, \\ & x_i(t) - x_j(t) \geq L, \quad \text{for } (i, j) \in \mathcal{C} \cap \mathcal{N}_r, \end{aligned}$$

where we again use subscript r to group variables according to their associated route.

2.2 Decomposition for delay objective

Suppose we use the crossing at the last intersection as performance measure, by defining the objective function as

$$J(x_i) = \inf\{t : x_i(t) \in \mathcal{E}_r(v_r(m_r))\}.$$

We show how to reduce the resulting problem to a scheduling problem, like we did in the single intersection case. It is not clear whether vehicles will always cross intersections at full




Figure 6: Empty disjunctive graph belonging to a tandem of two intersections, labeled $(1, 1)$ and $(2, 1)$. There are three routes $\mathcal{R} = \{0, 1, 2\}$. On route 0, four vehicle are arriving, the other two routes have 3 arrivals each. Conjunctive arcs are drawn as from left to right, travel arcs are drawn from top to bottom and disjunctive arcs are drawn as dotted lines. Two buffer constraints are drawn as diagonal arcs, corresponding to a capacity of 2 vehicles for the lane between both intersections.

speed, but we will simply require vehicles to do so from here on. Furthermore, we will again assume that all vehicles share the same geometry. Hence, the occupation time $\sigma \equiv \sigma(i, v)$ is the same for all vehicles and intersections. For this reason, we will write the shorthand $y_r \in \mathcal{S}_r$, because σ_r is no longer a free variable.

As a consequence of Assumption 1.1 and Assumption 1.2, each lower-level trajectory optimization problem for a given route $r \in \mathcal{R}$ decomposes into a sequence of problems, each corresponding to two consecutive intersection along V_r . This means that $y_r \in \mathcal{S}_r$ is equivalent to $y_{(v,w)} \in \mathcal{S}_{(v,w)}$ for each $(v, w) \in E_r$, where $y_{(v,w)}$ denotes the vector of all variables $y(i, v)$ and $y(i, w)$ for all $i \in \mathcal{N}_r$ and $\mathcal{S}_{(v,w)}$ denotes the set of values of $y_{(v,w)}$ for which a feasible trajectory part can be found. Hence, we will now focus on a tandem of two intersections and investigate the trajectories of vehicles in this with the goal of stating sufficient conditions for $y_{(v,w)} \in \mathcal{S}_{(v,w)}$.

2.3 Crossing time scheduling

- Introduce general job-shop problem and disjunctive graph.
- Introduce travel constraints and its disjunctive graph arcs.
- Introduce buffer constraints and its disjunctive graph arcs.
- Formulate MILP problem and investigate how the solving time scales with network size in terms of number of intersections and number of vehicles in the network. Do the single intersection cutting planes still hold? Are there any obvious cutting planes?

3 Constructive heuristics

Like for problems with a single intersection, we want to investigate possible constructive heuristics. In this case, the step-by-step construction can again be understood in terms of transitioning from a partial disjunctive graph to the next partial disjunctive graph by fixing

the direction of some disjunctive arc. Instead of only the route index, we now also need to specify the intersection to identify a unique disjunctive arc, hence the following definition.

Definition 3.1. Each pair (r, v) of a route $r \in \mathcal{R}$ and some intersection $v \in V_r$ which vehicles on r encounter is called a *crossing*. We say a crossing is done (pending) when all its vehicles are done (pending). Similarly, we say an intersection is done (pending) when all its crossings are done (pending). Therefore, the set of *valid actions* is exactly the set of pending crossings.

Similarly as in the single intersection scheduling problem, we will now consider heuristics that select the next action given the current partial disjunctive graph, which determines the next vehicle for which the crossing time is fixed. Before we turn our attention to deep learning models, we will consider a straightforward extension of the threshold heuristic.

3.1 Exhaustive heuristic

It is straightforward to extend the exhaustive policy (threshold heuristic with $\tau = 0$) to networks of intersections. When the network of intersections is a Directed Acyclic Graph (DAG), we can visit the intersections in topological order. When an intersection is visited, we apply the single intersection exhaustive scheduling policy on the disjunctive subgraph induced by the nodes belonging to that intersection in order to obtain the local ordering. We will use the performance of this heuristic as a baseline when evaluating heuristics based on a neural architecture. Furthermore, it also provides a *baseline* during reinforcement learning, reducing the variance of the REINFORCE estimator.

3.2 Neural constructive heuristic

We want to model optimal solutions using a similar auto-regressive model like we used for a single intersection. This means that we want to predict the next action, given the current state of the automaton, or equivalently, the current partial disjunctive graph. Before we consider models that can generalize to larger networks by using parameter sharing and message-passing schemes, we investigate the fit of a model that is based on recurrent embeddings of the horizons of crossings, similar to the recurrent embeddings of the horizons in the single intersection case.

We start by considering the imitation learning setting. Therefore, we sample a collection of training instances, which are solved to optimality using the MILP formulation. Next, we use the resulting crossing time schedule to compute actions for the automaton that lead to the same schedule. However, this sequence of actions is not unique: the order in which intersections are considered does not matter for the final schedule, see Figure 8. Instead of a single next action, the model should model the set of valid actions for every state of the automaton. These types of problems, where the output is required to be a set, are called multi-label classification.

Let us start by ignoring the multi-label nature of the problem by defining a model that selects a single next action. For each crossing, we define a *horizon* as a sequence of arrival time lower bounds as follows. Each horizon starts at the next unscheduled vehicle for each crossing. At each intersection, we compute the minimal crossing time lower bound across all unscheduled vehicles and we subtract this from all the lower bounds in each horizon at this intersection. Next, each horizon is embedded using an Elman RNN. These embeddings are then fed through a feedforward network to produce a logit for each crossing, see Figure 7.

Expert demonstration state-action pairs can be collected in different ways. For each solved instance in the training set, we can randomly select a single intersection order and replay the sequence of actions on the automaton to generate the corresponding sequence of state-action pairs. However, the model might become more robust when training on multiple samples of intersections orders per instance. Another possible strategy is to use a kind of fixed lookahead sampling procedure, as illustrated in Figure 9.




Figure 7: Illustration of model with RNN encoding of horizon at each crossing. The disjunctive graph nodes whose lower bounds are part of the current horizons are indicated with rectangles. The RNN embeddings (open circles) are fed through a final feedforward network to produce a logit (dots) for each crossing.

Instead of ignoring the multi-label nature of the problem, we can also design models that explicitly dictate the intersection visit order.

3.3 Network-agnostic architecture

We now focus on models that generalize across network sizes. In particular, we are interested in models that can be learned on relatively small networks (for which exact solutions might be tractable) and can then be applied to larger networks.

We use a GIN to compute an embedding for each node, which is then fed through an MLP and softmax to produce a probability over nodes. The GNN computes node embeddings, which are mapped to a score for each node. We compute the softmax over the scores of the nodes and then compute the negative log likelihood loss for backpropagation. Like in Zhang et al., each action corresponds to a unique node, encoding the operations that is dispatched next. However, we only really need to provide a route-intersection pair, but how to exploit this in the policy model? At this point, during the collection of observation-action pairs, we copy the whole disjunctive graph for each state. Alternatively, we could use some sort of masking for non-final states.

References

- [1] R. F. Hartl, S. P. Sethi, and R. G. Vickson, “A Survey of the Maximum Principles for Optimal Control Problems with State Constraints,” *SIAM Review*, vol. 37, no. 2, pp. 181–218, 1995.




Figure 8: When we fix the ordering of actions locally at every intersection, the global order of actions is not unique, because the local sequences may be merged in any way. Suppose we have a tandem of two intersections and the horizontal arrows correspond to taking the next local action at intersection 1, the vertical arrows correspond to taking next local action at intersection 2. The figure shows two possible global orders. Of course, the exact global order does not matter for the final schedule. However, it might be that trajectories near the boundary of the grid are harder to learn from data than trajectories that keep closer to the diagonal. The intuition is that we need to “look into the future” more to learn the former, while in the latter trajectories, progress in the two intersections is more balanced.




Figure 9: Possible strategy for collecting state-action pairs: one action look-ahead. At every state s with a set of possible actions $\mathcal{A}(s)$, we add all state-action pairs $\{(s, a) : a \in \mathcal{A}(s)\}$, but we pick a single action a^* to move the automaton further. In the figure, the state-action pairs are depicted as solid arrows. The encircled dots are the states that are actually visited. Observe that this procedure can be naturally generalized to a look-ahead of arbitrary depth.