
COORDINATION OF CONNECTED VEHICLES USING REINFORCEMENT LEARNING

Jeroen van Riel
Eindhoven University of Technology
j.a.a.v.riel@student.tue.nl

January 26, 2024

1 Motivation

Given the growing number of connected vehicles (CVs), an interesting question becomes how traffic can be handled efficiently in scenarios where every vehicle can be fully controlled by a single traffic manager with full information and perfect communication. Studying this question is important because it provides an upper bound on how efficiently traffic could be handled in such a scenario by neglecting issues related to communication latency, partial information, and decentralized control.

The problem we would like to consider is in essence about finding trajectories of vehicles through a network of intersections while optimizing for some measure of (economic or social) efficiency, to which we will refer as *trajectory planning*. Even when the route of a vehicle is fixed, there is still a lot of freedom in determining the speed profile over the course of this route. Previous studies have considered methods that are based on a two-stage decomposition: first determine the times when a vehicle crosses each intersection along its route, then compute the speed profile based on these times.

In general, efficiently planning the movement of vehicles through a road network is a difficult problem because control of one intersection directly influences the arrival process of vehicles of nearby intersections. In dense urban networks, these interactions become highly complex, which makes modeling the dynamics from first principles very difficult. Furthermore, the inherent nonstationarity of traffic demand, with fluctuations throughout the day, makes it difficult to design controllers that adapt their policy to these changes.

Deep learning methods have been proven to be very useful in capturing highly complex relationships that are otherwise very hard to capture from first principles. Furthermore, deep reinforcement learning methods have been successfully applied to automatically select good policies from very high-dimensional classes of policies, parameterized through deep neural networks. The aim of this proposal is to explore how deep reinforcement learning could be employed to develop procedures for trajectory planning of CVs in existing urban networks of intersections.

2 Literature Analysis

This section highlights some of the literature on traffic control by discussing illustrative classical methods for signalized intersections and giving an overview of recent applications of reinforcement learning techniques. The second part discusses relevant work on idealized traffic models featuring only connected autonomous vehicles (CVs).

2.1 Signalized Intersections

A vast amount of literature is available on methods for predicting and controlling road traffic. In particular, the context of networks of signalized intersections has received a substantial amount of attention. In most of these works, the main goal is to derive policies for setting traffic light signals such that traffic is handled as efficiently as possible, which is commonly measured in terms of the total delay experienced by all vehicles over some period of time. It has often been recognized that efficient control schemes should include some degree of coordination among the signal controllers at

individual intersections, which is commonly referred to as *coordination*. A well-known example of such a strategy is given by so-called *green waves* on arterial roads.

Early methods such as MAXBAND [1] aimed to synchronize multiple neighboring intersections, by tuning the timing offset between cyclic signal plans of the intersection, in order to achieve coordination. An interesting derivative of the original MAXBAND formulation is the PAMSCOD system [2], which is a framework that consists of a model for identifying platoons and a Mixed-Integer Linear Program (MILP) formulation for scheduling these platoons throughout a network in an online fashion. Under the assumption that the traffic controller receives actual position and speed information from a certain percentage of vehicles, a heuristic algorithm is proposed to identify platoons and estimate their size and expected arrival time at downstream intersections. The order in which these platoons cross the intersections in the network is then optimized using a MILP formulation that minimizes a combination of delay experienced at the next two encountered intersections. A solution provides the green times of signals for a fixed number of cycles in the future, so their method is a form of *rolling horizon optimization*.

Because deploying control policies in existing road infrastructure is often non-trivial, evaluation is often done by using a so-called *traffic microsimulator* like SUMO [3] or VISSIM. In contrast to *macroscopic models*, which aim to capture traffic dynamics in terms of aggregated quantities for large numbers of vehicles, such simulators try to capture the behavior of individual vehicles very precisely, such that realistic behavior on a larger scale emerges automatically.

Recently, there has been a growing interest in using reinforcement learning algorithms for developing efficient traffic signal control policies [4]. Most of this work is centered around a microsimulator to which an off-the-shelf reinforcement learning algorithm is applied in combination with some kind of neural function approximation to encode the high-dimensional state space. The obtained results are often very impressive in terms of pure performance, while relying on the availability of lots of computational power. Unfortunately, due to the model-free nature of the learning methods used, these studies do not provide much insight into the learned policies that could be exploited to derive methods that are computationally more efficient. Furthermore, the challenge of automatically learning how to achieve coordination among multiple intersections without any prior knowledge has been identified as one of the key remaining issues in this line of research.

2.2 Connected Vehicles

Assuming that all vehicles in the network are fully connected poses a lot of interesting new research questions [5]. However, the main problem of allocating intersection access time to vehicles remains a central issue. In this context, classical methods like MILP solving [6] are still an important tool for solving such *temporal* optimization problems. However, the current setting allows us to consider optimization in the *spatio-temporal* domain. With signalized intersections and human drivers, trajectories of individual vehicles cannot be directly controlled, but under the current assumption *trajectory planning* becomes relevant. This issue has already been studied for the *conflict zone* of intersections. For example, in [7] the trajectories of vehicles are modeled as volumes in time and space that can be used to characterize conflicting trajectories.

For a single intersection, the two-stage approach in [8] first computes access times using dynamic programming and use these to compute conflict-free trajectories as solutions to an optimal control problem. A similar approach was taken in [9], where intersection access times are determined by simple policies from queueing theory instead. The performance of their method was analyzed using results from *polling theory*.

When the optimization objective only considers delay at the intersection, i.e., only in the temporal domain, the problem of finding efficient access times can be interpreted as a variant of *single machine scheduling*. More precisely, using terminology from the *machine scheduling* literature [10], it can be shown that the problem reduces to a single machine scheduling problem with release dates, job families with chain precedence constraints (corresponding to lanes), family-dependent setup times and total completion time objective.

Machine scheduling is a widely studied subfield of combinatorial optimization with a wide range of applications in, for example, manufacturing and healthcare. Dealing with situations in which not all information about jobs is available upfront is known as *online scheduling* [11, 12]. The performance of an online scheduling policy depends highly on how information is disclosed, so different metrics are necessary for evaluation. A common approach is to consider the worst-case performance compared an optimal *offline algorithm* that has access to all future information. Alternatively, an explicit specification of the information disclosure process may be assumed to analyze performance in expectation. This kind of analysis is particularly important for our current context of road traffic control, because demand for mobility is highly unpredictable.




Figure 1: Example graph G with three intersection i_1, i_2, i_3 and three vehicles routes, indicated by R_1, R_2 and R_3 . The external nodes are labeled as e_1, \dots, e_6 . The weight (travel time) of each arc is indicated next to it.

3 Problem Formulation and Objective

This section provides a sketch of our model of coordination of connected vehicles through a network of intersections which will serve as the main conceptual model throughout the project. Vehicles are represented as little dots and arrive to a graph over time. They follow a fixed deterministic route towards a final node, where the vehicle leaves the system. The goal of the traffic controller is to determine how vehicles move exactly along their route by providing their speed profiles, while satisfying some constraints on the interaction between vehicles in lanes and at intersections that are due to safety considerations.

Let G be some road network, modeled as a directed graph, with intersections represented by nodes and lanes represented by arcs. We assume that vehicle are not allowed to overtake, so they need to stay in their lane. We assume that routes are fixed and deterministic. For each vehicle j , let its route be denoted by R_j , which is encoded as a sequence of nodes. Vehicle j enters the network at some external node $R_j(0)$ at time r_j and then visits the n_j intersections $R_j(1), \dots, R_j(n_j)$ until it reaches the exitpoint $R_j(n_j + 1)$, where it leaves the network. See Figure 1 for an illustration.

We assume that there is one *central traffic controller* that knows the exact position of every vehicle in the network. Let $x_j(t)$ denote the position of vehicle j on its route R_j at time $t \geq r_j$. We also use the notation $x_{ij}(t)$ to denote the position in the lane on the route of j leading to the i th intersection on its route, see Figure 2. Furthermore, let $v_j(t)$ and $a_j(t)$ denote the speed and acceleration, respectively, of vehicle j at time t .

The traffic controller determines trajectories $x_j(t)$ satisfying some kind of smoothness and safety constraints. For example, we may impose maximum/minimum speed and acceleration constraints

$$v_{\min} \leq v_j(t) \leq v_{\max}, \quad (1)$$

$$a_{\min} \leq a_j(t) \leq a_{\max}. \quad (2)$$




Figure 2: Illustration of a network of intersections with the position of some vehicle j given in two ways.

Given a node i , let $t_{ij} = t_{ij}(x_j)$ denote the time when vehicle j crosses i . Between crossing of vehicles from different lanes, we enforce a *safe lane-switching* time s_{lane} . Whenever a pair of distinct vehicles $j \neq l$ crosses the same intersection i , we require that either

$$t_{ij} + s_{\text{lane}} \leq t_{il} \quad \text{or} \quad t_{il} + s_{\text{lane}} \leq t_{ij}. \quad (3)$$

Furthermore, when two vehicles j and l are driving in the same lane we require some minimum *safe following distance* d_{follow} . For all pairs of distinct vehicles $j \neq l$, we require that

$$x_{ij}(t) + d_{\text{follow}} \leq x_{il}(t) \quad \text{or} \quad x_{il}(t) + d_{\text{follow}} \leq x_{ij}(t), \quad (4)$$

at all times t when both vehicles are driving towards a common intersection i .

The objective for the traffic controller is to compute trajectories while minimizing some performance measure, e.g., the total delay experienced by vehicles at intersections, which is equivalent to minimizing

$$\sum_j \sum_{i \in R_j} t_{ij}, \quad (5)$$

where the sum is taken over all indices of vehicles that were present in the network during a certain time interval.

3.1 Objective

In general, our goal is to develop computationally efficient methods for solving the problem sketched above. More specifically, we aim to employ existing off-the-shelf reinforcement learning techniques to learn policies for efficient trajectory planning.

A proven approach is to decompose the solution into a part that schedules the intersection access times and a part that uses this schedule to compute continuous conflict-free trajectories. Inspired by the idea of integrating machine learning methods with combinatorial optimization [13], we aim to study an alternative method of scheduling access times based on a graph representation that allows us to apply reinforcement learning [14].

Besides exploiting this explicit two-stage decomposition, we would like to consider a spatio-temporal scheduling model to generate trajectories in a more direct way. We have a preliminary solution in this direction that we would like to validate. Ideally, we would also like to learn policies for this model using reinforcement learning.

4 Research Approach

This section provides an overview of our current concrete research ideas. Because temporal scheduling is still not fully understood in even very simple cases, we will at first ignore the spatial issues by focusing on methods of intersection access scheduling for single intersections and networks of intersections. After that, we propose a model for developing trajectory planning algorithms that do not rely on a two-stage scheduling decomposition.

4.1 Single Isolated Intersection

Before investigating coordination across multiple intersections, we first revisit the special case with a single isolated intersection. This already offers many interesting questions that may help us later in distinguishing which problems arise due to network effects or spatial constraints.

The single intersection scheduling problem can be thought of as a variant of a single machine scheduling problem, in which jobs need to be assigned time on a single machine. Like other single machine problems, the current problem can be rather naturally formulated as a Mixed-Integer Linear Program (MILP) by introducing binary decision variables for the ordering of vehicles from different lanes. Relying on the current power of MILP solvers, e.g. Gurobi [15], this provides an exact solution to the problem, assuming that all arrivals are known from the start.

In order to apply reinforcement learning, we need to precisely specify the sequential decision making process as a Markov Decision Process (MDP). We propose to study a simple *dispatching* policy that computes a schedule by simulating the behavior of a traditional traffic light. More precisely, the policy decides which vehicle is allowed to cross the intersection next in the following way. The agent takes a vector of the next h earliest crossing times for both lanes as input and decides whether to keep serving the current lane or to switch to the other lane. The number h of next crossing times the controller sees is referred to as the *horizon*. In a sense, this resembles the way traffic lights switch between *phases* at most existing signalized intersection.

Choosing the most efficient reinforcement learning algorithm for our current model is not our main concern. Therefore, we will settle with an off-the-shelf implementation of deep Q-learning with experience replay [16] for this single intersection environment. Once a policy has been trained that performs reasonably well in terms of the optimization objective, we can analyze it in the following ways:

- **Comparison.** First of all, we can compute optimal solutions for small instances, so the approximation ratio can be computed precisely in these cases by comparing the objective value to the optimal found by a MILP solver. Furthermore, the performance can be compared to existing policies under the assumption of a fixed distribution of the arrivals r_j . Examples of policies to compare with are exhaustive service, in which the current lane is served as long as vehicles are waiting there, or the gated policy from [9].
- **Effect of horizon size.** It is possible to experiment with different sizes of the horizon. For a fixed distribution of the arrivals r_j , we expect that the performance of the learned policy does not improve significantly beyond a certain horizon size.
- **Generalization.** A central question in reinforcement learning is the issue of generalization. For example, we could assess how well the trained policy performs with a (slightly) different distribution of r_j . Furthermore, when considering more than two incoming lanes, an interesting question is whether we can train policies that generalize across different numbers of lanes.
- **Policy interpretation.** For learned policies that perform really well, it might be worthwhile to analyze their behavior a bit further in the hope of discovering structure that can be used to define better policy spaces. Two possible ways of doing this are by providing certain “test states” to see which actions get selected and providing test instances for which we know the optimal solution to see how close the schedule produced by the policy is.

4.2 Network of Intersections

The next step is to consider networks of connected intersections. Before considering the problem in the spatial and temporal domain in a unified fashion, we treat the scheduling problem for multiple intersections. More precisely, we again assume there is no limit on the number of vehicles that can be present in lanes between intersection. This provides a lot of new interesting questions, because decisions taken at an intersection influence the arrival process of downstream intersections. Under these assumptions, we obtain a variant of the well-known and widely studied job shop scheduling problem [10]. Job shop problems are a particular type of multi-machine scheduling problem in which jobs consist of consecutive stages, where each stage must be executed on a different machine.

Again, the offline variant of this problem can be formulated as a MILP and solved to optimality with readily available solvers. However, we expect that this approach does not scale well in terms of the network size and number of vehicles. Therefore, in order to solve problems at a real-world scale, we need to focus on finding good heuristics. Our results for the single isolated intersection can be readily applied to the current model by controlling each intersection independently by the single intersection dispatching policy. While this is still interesting to try, a different approach is required in order to obtain coordination among intersections.

Instances of job shop scheduling can be adequately represented as a graph that encodes the precedence constraints between jobs. This graph is called the *disjunctive graph*, because a subset of the arcs encodes the disjunctive ordering decisions between jobs on the same machine (vehicles approaching the same intersection in the current context). A recent study [14] has exploited this structure to obtain a principled approach to applying reinforcement learning to the job shop problem. Their current method is only applicable as an offline scheduling algorithm, so we would like to extend their method such that it can also be applied in an online setting.

In addition to the questions discussed in the previous section, the current network setting provides the following additional directions of inquiry.

- **Coordination along arterial road.** A widely studied phenomenon in traffic signal control is coordination along a series of arterial intersections. The idea is that aligning the timing of traffic signal phases along a series of connected intersections that handle relatively large fraction of total traffic is often a good strategy to optimize overall delay in the network. Consider a network with a couple of intersections in series with some minor adjacent roads. By using an instance with a very regular arrival pattern on the main arterial and a limited arrival rate at the minor roads, we can assess coordination by measuring whether the policy actually exploits the regularity of the arrivals.
- **Platoon splitting.** For a single intersection, it has been shown that splitting platoons of vehicles is never beneficial [17]. However, for more than one intersection, it can be shown that this platoon preservation theorem

no longer holds, even in some very simple situations. It would be interesting to understand better why this happens by characterizing this kind of situation.

- **Generalization across network topologies.** Ideally, the scheduling method can be used regardless of the network topology at hand. In that case, it is possible to study how well a policy trained on some network topology generalizes to other topologies. This further motivates our idea of using the job shop disjunctive graph, which could help in making the policy topology-agnostic.

4.3 Finite Buffers Model

In the models discussed above, we assumed that vehicles do not interact in lanes, hence it is possible for an unbounded number of vehicles to wait in a lane at any given time. However, in real-world traffic networks, the maximum possible number of vehicles is limited by spatial constraints and even depends non-trivially on the speeds of individual vehicles. Therefore, we propose a model that incorporates this aspect by considering the position of vehicles in lanes. Instead of encoding the precise location in a continuous sense, we divide each lane in a finite number of locations. For a given vehicle, the controller has to decide how long it should wait on each location. For large number of locations per lane, this method allows us to approximate continuous trajectories through interpolation.

We illustrate our idea for a single intersection. Assume that it takes constant time Δt for a vehicle to travel to the next location. For each lane k , let m_k denote the number of locations, excluding the entrypoint. Let the locations of lane k be identified by increasing numbers $\mathcal{L}(k) = (1, \dots, m_k)$, where the last one corresponds to the intersection. In the following variable definitions, we will not explicitly mention the dependency on the lane to keep notation simple. Let y_{ij} denote the time at which vehicle j departs from location $i \in \{0, \dots, m_k\}$, then $y_{ij} + \Delta t$ is the arrival time of vehicle j at location $i + 1$. To simplify notation, we define $\bar{y}_{ij} = y_{i-1,j} + \Delta t$ to be the arrival time of vehicle j at location i . For every location i and vehicle j , we require

$$\bar{y}_{ij} \leq y_{ij} \quad (6)$$

and the difference between these two quantities is called the *location delay*. For each pair of consecutive vehicles in the same lane k with precedence constraint $j \rightarrow l$ (so vehicle l arrived later than j), we have the constraints

$$y_{ij} + p \leq \bar{y}_{il}, \quad (7)$$

for every location i in their lane. Next, we model the safety constraints involving vehicles from different lanes that cross a common intersection i . Let j be some vehicle in lane $k(j)$, then let y_j denote the departure time from the intersection, so we have $y_j = y_{ij}$ with $i = m_{k(j)}$. Similarly, let \bar{y}_j denote the arrival time of j at the intersection, so we have $\bar{y}_j = y_{i-1,j} + \Delta t$ with $i = m_{k(j)}$. From constraints (7), we see that it makes sense to say that vehicle j *occupies the intersection* during the interval

$$[\bar{y}_j, y_j + p]. \quad (8)$$

We require an additional *switch-over time* s whenever the next vehicle to occupy the intersection comes from a different lane. This results in the additional constraints

$$y_j + p + s \leq \bar{y}_l \quad \text{or} \quad y_l + p + s \leq \bar{y}_j, \quad (9)$$

for all pairs of conflicting vehicles $j \neq l$ that approach intersection i from distinct lanes.

Because of the discrete nature of the model, it is straightforward to formulate it as a MILP, which allows us to obtain optimal schedules y for objectives similar to (5). However, as the number of variables grows rapidly with the values of m_k , we do not expect this approach to scale well. Therefore, we would like to investigate how we could employ reinforcement learning to solve this problem, similarly to how we do this for the job shop variant for network scheduling.

A major question is how to structure the policy for determining location delays for vehicles in a online setting. It might seem straightforward to let the controller determine, for every vehicle, the location delays of the next couple of locations. An alternative approach would be to have the controller set a location delay for each location, independently of the vehicles. In a sense, this is somewhat similar to the dispatching policy for the single intersection scheduling problem, in that the problem is agnostic with respect to the number of vehicles present in the network.

5 Timeline

This section provides a rough sketch of the timeline of the project from now on. Counting from the week of February 5th onwards, I propose a rough division of the project over approximately 22 weeks, aiming to finish by the end of July 2024.

At this point, we have already developed some methods for each of the three models. We are currently analyzing the trained policies for the single intersection case. After this, we try to extend the job shop scheduling method with disjunctive graphs to online problems. The ordering of parts in the following plan is based on the decomposition above. However, we are considering working on the interpolation method for the finite buffers model when we are finished with the analysis of the single intersection, because we want to have an early proof of concept to check whether the proposed approach actually makes sense.

- **Single intersection environment (done).** Develop a Gymnasium [18] environment for the single intersection dispatching policy.
- **DQN agent for single intersection (done).** Using the off-the-shelf DQN implementation provided by CleanRL [19], we managed to train good policies for some examples of arrival distributions.
- **Analyze performance of single intersection policy (started, 2 weeks).** We are currently comparing the performance of the trained policies with optimal solutions found by solving the corresponding MILP for small instances. As we indicated in Section 4.1, we would also like to study the effect of the horizon size h and some basic study of generalization over different arrival distributions.
- **Interpret single intersection policy (2 weeks).** We think that there is some particular structure to the optimal policies for the simple single intersection scheduling problem. Therefore, we would like to try understanding the behavior of the trained policies a bit better in the hope of discovering some simple rules that can be used to develop more efficient algorithms.
- **Adapt disjunctive graph-based scheduling (4 weeks).** The job shop scheduling method based on the disjunctive graph [14] needs some slight adaptations in order to be applicable to the variant of job shop that is relevant for scheduling in networks. The authors provide reinforcement learning code and a Gym environment. However, their code is not very well documented, so we are considering a rewrite of the necessary parts.
- **Analyze network policies (1 week).** Once we are able to solve the job shop variant in an offline setting, we can do some preliminary experimentation and performance analysis.
- **Define how to adapt the disjunctive graph for online scheduling (2 week).** We have a rough idea of how the disjunctive graph scheduling method could be extended to be applicable in an online fashion. The basic idea is to just extend the disjunctive graph every time new vehicles (jobs) arrive. Therefore, the main challenge is finding a suitable deep learning embedding of this graph that can deal with these kinds of extensions.
- **Implement online scheduling (3 weeks).** We expect that we need to reconsider the graph neural network that is used to learn an embedding of the disjunctive graph in order to get our intended method working.
- **Analyze online performance of network policies (2 weeks).** Like we did for the single intersection case, the learned policies should be analyzed in terms of performance. Furthermore, we should investigate whether the policies show some degree of coordination on intersections along arterial roads.
- **Define interpolation for finite buffers model (2 weeks).** Our current definition of the finite buffers model does not yet involve continuous trajectories. We still need to exactly define how the location delays can be used to compute non-overlapping vehicle trajectories.
- **Finite buffers model environment (started, 2 weeks).** Given the interpolation method, we are ready to develop a Gymnasium [18] environment of the finite buffers model.
- **Prototype reinforcement learning for finite buffers (2 weeks).** Develop a working reinforcement learning procedure based on policies that set the location delays in a vehicle-agnostic fashion.

References

- [1] William R. McShane and Roger P. Roess. *Traffic engineering*. Prentice Hall polytechnic series in traffic engineering. Prentice-Hall, Englewood Cliffs, N.J., 1990.
- [2] Qing He, K. Larry Head, and Jun Ding. PAMSCOD: Platoon-based arterial multi-modal signal control with online data. *Transportation Research Part C: Emerging Technologies*, 20(1):164–184, February 2012.
- [3] Pablo Alvarez Lopez, Evamarie Wiessner, Michael Behrisch, Laura Bieker-Walz, Jakob Erdmann, Yun-Pang Flotterod, Robert Hilbrich, Leonhard Lucken, Johannes Rummel, and Peter Wagner. Microscopic Traffic Simulation using SUMO. In *2018 21st International Conference on Intelligent Transportation Systems (ITSC)*, pages 2575–2582, Maui, HI, November 2018. IEEE.
- [4] Mohammad Noeen, Atharva Naik, Liana Goodman, Jared Crebo, Taimoor Abrar, Zahra Shakeri Hossein Abad, Ana L.C. Bazzan, and Behrouz Far. Reinforcement learning in urban network traffic signal control: A systematic literature review. *Expert Systems with Applications*, 199:116830, August 2022.

- [5] Ashkan Gholamhosseinian and Jochen Seitz. A Comprehensive Survey on Cooperative Intersection Management for Heterogeneous Connected Vehicles | IEEE Journals & Magazine | IEEE Xplore. *IEEE Access*, Volume 10.
- [6] Seyed Alireza Fayazi and Ardalan Vahidi. Mixed-Integer Linear Programming for Optimal Scheduling of Autonomous Vehicle Intersection Crossing. *IEEE Transactions on Intelligent Vehicles*, 3(3):287–299, September 2018. Conference Name: IEEE Transactions on Intelligent Vehicles.
- [7] Zhenning Li, Qiong Wu, Hao Yu, Cong Chen, Guohui Zhang, Zong Z. Tian, and Panos D. Prevedouros. Temporal-spatial dimension extension-based intersection control formulation for connected and autonomous vehicle systems. *Transportation Research Part C: Emerging Technologies*, 104:234–248, July 2019.
- [8] Yiheng Feng, Chunhui Yu, and Henry X. Liu. Spatiotemporal intersection control in a connected and automated vehicle environment. *Transportation Research Part C: Emerging Technologies*, 89:364–383, April 2018.
- [9] R. W. Timmerman and M. A. A. Boon. Platoon forming algorithms for intelligent street intersections. *Transport-metrica A: Transport Science*, 17(3):278–307, February 2021.
- [10] Michael L. Pinedo. *Scheduling: Theory, Algorithms, and Systems*. Springer International Publishing, Cham, 2016.
- [11] Rita Guerreiro, André S. Santos, and Anabela Tereso. Online Scheduling: A Survey. In *2023 18th Iberian Conference on Information Systems and Technologies (CISTI)*, pages 1–6, June 2023. ISSN: 2166-0727.
- [12] A.P.A. Vestjens. *On-line machine scheduling*. Phd Thesis 1 (Research TU/e / Graduation TU/e), Technische Universiteit Eindhoven, Eindhoven, 1997. ISBN: 9789038605715.
- [13] Yoshua Bengio, Andrea Lodi, and Antoine Prouvost. Machine Learning for Combinatorial Optimization: a Methodological Tour d’Horizon, March 2020. arXiv:1811.06128 [cs, stat].
- [14] Cong Zhang, Wen Song, Zhiguang Cao, Jie Zhang, Puay Siew Tan, and Chi Xu. Learning to Dispatch for Job Shop Scheduling via Deep Reinforcement Learning, October 2020. arXiv:2010.12367 [cs, stat].
- [15] Gurobi Optimization, LLC. Gurobi optimizer reference manual, 2024.
- [16] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andrei A. Rusu, Joel Veness, Marc G. Bellemare, Alex Graves, Martin Riedmiller, Andreas K. Fidjeland, Georg Ostrovski, Stig Petersen, Charles Beattie, Amir Sadik, Ioannis Antonoglou, Helen King, Dhharshan Kumaran, Daan Wierstra, Shane Legg, and Demis Hassabis. Human-level control through deep reinforcement learning. *Nature*, 518(7540):529–533, February 2015. Number: 7540 Publisher: Nature Publishing Group.
- [17] Matthijs Limpens. *Online Platoon Forming Algorithms for automated vehicles: A more efficient approach*. Bachelor, Eindhoven University of Technology, September 2023.
- [18] Mark Towers, Jordan K. Terry, Ariel Kwiatkowski, John U. Balis, Gianluca de Cola, Tristan Deleu, Manuel Goulão, Andreas Kallinteris, Arjun KG, Markus Krimmel, Rodrigo Perez-Vicente, Andrea Pierré, Sander Schulhoff, Jun Jet Tai, Andrew Tan Jin Shen, and Omar G. Younis. Gymnasium, March 2023.
- [19] Shengyi Huang, Rousslan Fernand Julien Dossa, Chang Ye, Jeff Braga, Dipam Chakraborty, Kinal Mehta, and João G.M. Araújo. CleanRL: High-quality single-file implementations of deep reinforcement learning algorithms. *Journal of Machine Learning Research*, 23(274):1–18, 2022.