


Coordination of autonomous vehicles

Jeroen van Riel


January 2025

Coordination of autonomous vehicles




- Less human intervention (less traffic lights required)
- Better guarantees on safety
- Potentially reduce economic costs

... as optimal control problem



- Multi-agent optimal control problem
 - Minimize total travel time
 - Avoid collisions

...as optimal control problem



- Some simplifying assumptions
 - Central control with perfect communication
 - Fixed routes
 - All future arrivals known

How to solve it?

- Direct transcription methods
 - Provide optimal trajectories
 - Computationally very demanding
- How to solve large instances?
 - Need for approximation
 - Exploit problem structure (decomposition)
 - Automatically derive heuristics (learning)

Research questions

Q1: Exploit problem structure (decomposition)


How can we model the offline trajectory optimization problem, modeling coordination of autonomous vehicles in networks of unsignalized intersections, as a variant of job-shop scheduling to effectively reduce the dimensionality of the problem, while guaranteeing the generation of collision-free trajectories?

Q2: Automatically derive heuristics (learning)


Can recent deep reinforcement learning formulations for combinatorial optimization problems, like job-shop scheduling, be adapted to the offline trajectory optimization problem, formulated in terms of job-shop scheduling, in order to provide a scalable optimization algorithm that automatically learns to exploit hidden structures in problem instances?

Decomposition



- Upper-level crossing time scheduling
 - Mixed-Integer Linear Programming (MILP)




- Lower-level trajectory optimization problem
 - Direct transcription \rightarrow linear programming



Determine crossing times




Determine crossing order




- Crossing times follow from crossing order

Determine crossing order




- Map instance to optimal crossing order
- Use step-by-step construction...


Determine crossing order




Determine crossing order




Determine crossing order



Determine crossing order



Determine crossing order




Learn crossing order

- ~~Map instance to optimal crossing order~~
- Map partial order to next partial order (policy)
- We can learn this policy from examples!
 - Imitation learning from optimal MILP solutions
 - Reinforcement learning with dense delay reward

Overview of project plan



- Coordination as optimal control problem
- Vehicle scheduling + trajectory optimization
- Sequentially construct crossing order
- Learn from examples



Appendix: Disjunctive graph



Disjunctive graph

- Partial solutions encoded as disjunctive graph augmented with lower bounds on crossing times
- Parameterize ordering policy based on graph neural network embedding of augmented disjunctive graph



Disjunctive graph

- Partial solutions encoded as disjunctive graph augmented with lower bounds on crossing times - Parameterize ordering policy based on graph neural network embedding of augmented disjunctive graph




Appendix: Related literature

- Autonomous intersections
- Neural combinatorial optimization


Autonomous intersections

- “Autonomous Intersection Control” (Dresner & Stone)
 - Single intersection
 - Time slot reservation-based protocol
 - Central intersection manager




Autonomous intersections

- “Approximate Optimal Coordination” (Hult et al.)
 - Single intersection
 - Single vehicle per lane
 - Explicit collision-avoidance constraints




Neural combinatorial optimization

- “Learn to dispatch” (Zhang et al.)
 - Job-shop scheduling problem
 - Dispatch next operation
 - Policy using Graph Isomorphism Network (GIN)






Appendix: Single intersection



- Notation
- Upper-level crossing time scheduling
- Lower bound on starting times
- Imitation learning with neural policy
- Lower-level trajectory optimization

Notation

- vehicle indices \mathcal{N}
- $y(i)$ is crossing time of vehicle i
- r_i earliest crossing time of vehicle i



- i and j same lane: $y(i) + \rho \leq y(j)$
- i and j distinct lanes: $y(i) + \sigma \leq y(j)$ or $y(j) + \sigma \leq y(i)$

Upper-level crossing time scheduling

- conjunctive constraints \mathcal{C}
- disjunctive (conflict) constraints \mathcal{D}

$$\begin{array}{ll}\min_y & \sum_{i \in \mathcal{N}} y(i) \\ \text{s.t.} & r_i \leq y(i), \quad \text{for all } i \in \mathcal{N}, \\ & y(i) + \rho \leq y(j), \quad \text{for all } (i, j) \in \mathcal{C}, \\ & y(i) + \sigma \leq y(j) \text{ or } y(j) + \sigma \leq y(i), \quad \text{for all } (i, j) \in \mathcal{D}\end{array}$$

Upper-level crossing time scheduling

- Formulate as mixed-integer linear program (MILP)
- Introduce binary decision variables γ_{ij}
- Use big-M technique


$$\begin{array}{ll}\min_y & \sum_{i \in \mathcal{N}} y_i \\ \text{s.t.} & r_i \leq y_i, \quad \text{for all } i \in \mathcal{N}, \\ & y_i + \rho_i \leq y_j, \quad \text{for all } (i, j) \in \mathcal{C}, \\ & y_i + \sigma_i \leq y_j + \gamma_{ij}M, \quad \text{for all } (i, j) \in \mathcal{D}, \\ & y_j + \sigma_j \leq y_i + (1 - \gamma_{ij})M, \quad \text{for all } (i, j) \in \mathcal{D}, \\ & \gamma_{ij} \in \{0, 1\}, \quad \text{for all } (i, j) \in \mathcal{D}\end{array}$$

Lower bounds on starting times

- Disjunctive graph given current order π
- Nodes are vehicle indices \mathcal{N}
- Edges $i \xrightarrow{w(i,j)} j$
 - Conjunctive edges $i \xrightarrow{\rho} j$
 - Disjunctive edges $i \xrightarrow{\sigma} j$ or $j \xrightarrow{\sigma} i$
- Lower bounds LB_{π} on starting times given current order π

$$\text{LB}_{\pi}(j) = \max\{r_j, \text{LB}_{\pi}(i) + w(i,j)\}$$

Imitation learning with neural policy



- crossing order $\pi = ((1, 1), (2, 1))$ of vehicles
- step-by-step construction of this order
 - 1. choose $(1, 1)$
 - 2. choose $(2, 1)$
 - 3. ...

Imitation learning with neural policy

- get optimal trajectories from MILP solver
- parameterize policy based on LB_{π}
 - only consider $LB_{\pi}(j)$ for unscheduled j
 - recurrent embedding of $LB_{\pi}(j)$ per lane
 - alternatively, use zero padding
- fit policy parameters to expert transitions

Lower-level trajectory optimization

- position x , velocity v , control input u
- position of vehicle in front x' , follow distance L
- position of intersection B , crossing time τ

$$\begin{aligned} \arg \min_{x:[0,\tau] \rightarrow \mathbb{R}} & \int_0^\tau |x(t)| dt \\ \text{s.t. } & \ddot{x}(t) = u(t), & \text{for all } t \in [0, \tau], \\ & |u(t)| \leq a_{\max}, & \text{for all } t \in [0, \tau], \\ & 0 \leq \dot{x}(t) \leq v_{\max}, & \text{for all } t \in [0, \tau], \\ & x'(t) - x(t) \geq L, & \text{for all } t \in [0, \tau], \\ & (x(0), \dot{x}(0)) = s_0, \\ & (x(\tau), \dot{x}(\tau)) = (B, v_{\max}) \end{aligned}$$